

Gromit

An In-Memory Graph Database

Yunling Cui



University of Waterloo
yunling.cui@uwaterloo.ca

February 10, 2017

Overview

- 1 Motivation
- 2 Main Contributions
- 3 Gromit
 - Graph Storage
 - Query Framework
 - Transaction Management
 - Benchmarks
- 4 Summary

Social Network Applications



[3]

Relational Database Management System (RDBMS)

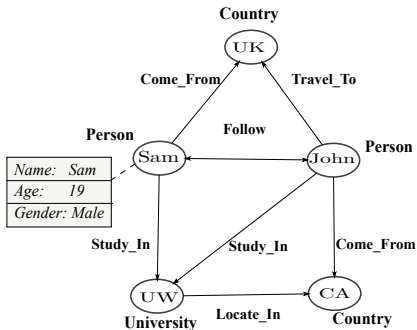


Table: Person			
P.ID	Name	U.ID	C.ID
0	Sam	2	4
1	John	2	3

Table: University			
U.ID	Name	Abbr.	C.ID
2	University of Waterloo	UW	3
5	University of Toronto	UofT	3

Table: Country		
C.ID	Name	Capital
3	Canada	Ottawa
4	UK	London

Relational Database Management System (RDBMS)

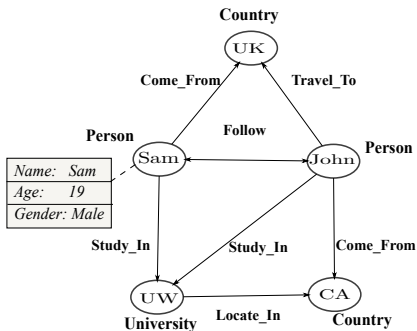


Table: Person			
P.ID	Name	U.ID	C.ID
0	Sam	2	4
1	John	2	3

Table: University			
U.ID	Name	Abbr.	C.ID
2	University of Waterloo	UW	3
5	University of Toronto	UofT	3

Table: Country		
C.ID	Name	Capital
3	Canada	Ottawa
4	UK	London

NoSQL Store

Examples

- Wide-Column Store [1]
- Key-Value Store [4]
- Document Store [2]

Person Document

ID	0
Name	Sam
University	2
Country	4

Person Document

ID	1
Name	John
University	2
Country	3

University Document

ID	2
Name	University of Waterloo
Abbr.	UW
Country	3

Country Document

ID	3
Name	Canada
Capital	Ottawa

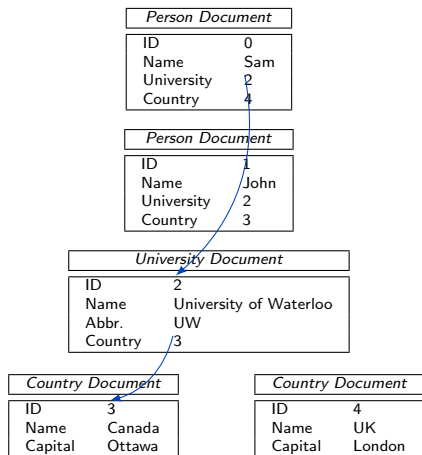
Country Document

ID	4
Name	UK
Capital	London

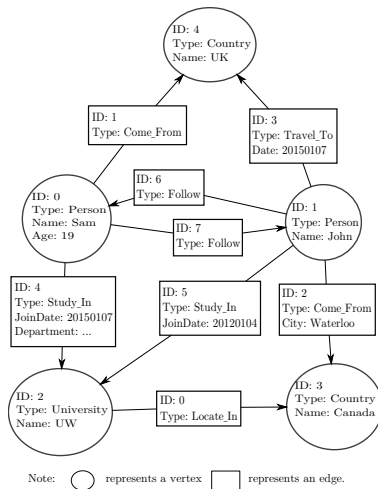
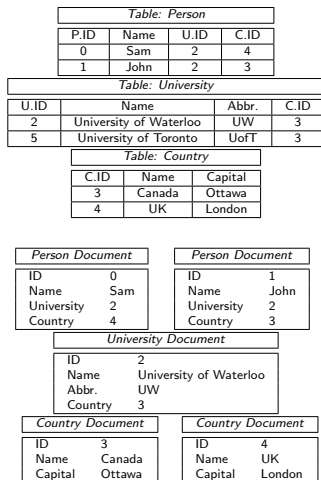
NoSQL Store

Examples

- Wide-Column Store [1]
- Key-Value Store [4]
- Document Store [2]



Graph Store



Motivation

- Graph databases are suitable for highly connected data
- Simulators work well with applications written in C++



Motivation

- Graph databases are suitable for highly connected data
- Simulators work well with applications written in C++



Objective: To build a graph database backend in C++

1 Motivation

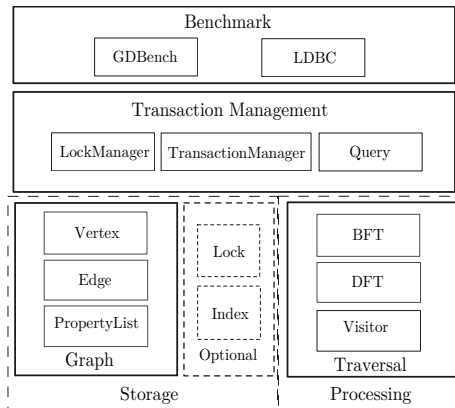
2 Main Contributions

3 Gromit

- Graph Storage
- Query Framework
- Transaction Management
- Benchmarks

4 Summary

Main Contributions



1 Motivation

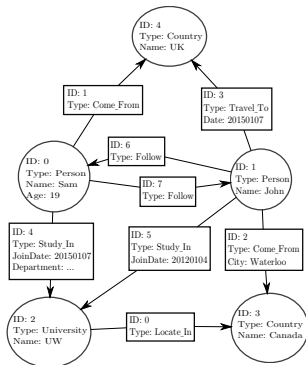
2 Main Contributions

3 Gromit

- Graph Storage
- Query Framework
- Transaction Management
- Benchmarks

4 Summary

Graph Storage



Vertex			
ID	Label	PropertyList	NextEdge
0	Person	Pv0	E7
1	Person	Pv1	E7
...			
4	Country	Pv4	E3
...			

PropertyList	
Key	Value
First Name	Sam
Last Name	Smith
...	
Date of Birth	...
...	

PropertyList	
Key	Value
Name	UK
Capital City	London
...	

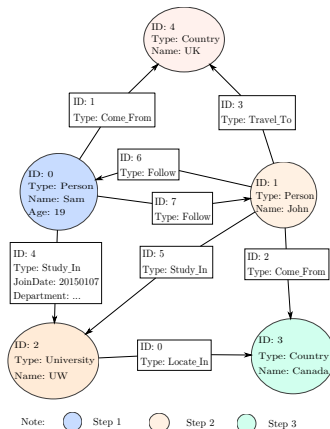
Edge								
ID	Label	PropertyList	FirstVertex	SecondVertex	FNE	FPE	SNE	SPE
...								
1	Follow	Pe1	V0	V4	E4	E7		E3
3	Travel_To	Pe1	V1	V4	E2	E7	E1	
...								
7	Follow	Pe1	V0	V1	E1		E3	
...								

Graph Processing

- Breadth-First Search (BFS)
- Depth-First Search (DFS)

Example

```
function BFS(Graph, Node)
  create empty set S
  create empty Queue Q
  Q.enqueue(Node)
  S.insert(Node)
  while !Q.empty() do
    Current = Q.dequeue()
    for all n in Current.neighbor do
      if !S.has(n) then
        S.insert(n)
        Q.enqueue(n)
      end if
    end for
  end while
end function
```



Query Framework with Visitor

```
function BFT(Graph, Node, Visitor)
```

```
    create empty set S
```

```
    create empty Queue Q
```

```
    Q.enqueue(Node)
```

```
    S.insert(Node)
```

```
    while !Q.empty() do
```

```
        Current = Q.dequeue()
```

```
        for all n in Current.neighbor do
```

```
            if !S.has(n) then
```


Query Framework with Visitor

```
function BFT(Graph, Node, Visitor)
  create empty set S
  create empty Queue Q
  Q.enqueue(Node)
  S.insert(Node)
  while !Q.empty() do
    Current = Q.dequeue()
    for all n in Current.neighbor do
      if !S.has(n) then
```

Filter 1	filterLabel(" University"); filterDirection(OUT);
Filter 2	filterDepth(1)

Query Framework with Visitor

```
function BFT(Graph, Node, Visitor)
  create empty set S
  create empty Queue Q
  Q.enqueue(Node)
  S.insert(Node)
  while !Q.empty() do
    Current = Q.dequeue()
    for all n in Current.neighbor do
      if !S.has(n) then
        if Visitor.visitNext(n) then
          return
        end if
        S.insert(n)
        Q.enqueue(n)
      end if
    end for
  end while
end function
```

Query Framework with Visitor

```
function BFT(Graph, Node, Visitor)
  create empty set S
  create empty Queue Q
  Q.enqueue(Node)
  S.insert(Node)
  while !Q.empty() do
    Current = Q.dequeue()
    for all n in Current.neighbor do
      if !S.has(n) then
        if Visitor.visitNext(n) then
          return
        end if
        S.insert(n)
        Q.enqueue(n)
      end if
    end for
  end while
end function
```

- Selection
- Summarization
- Path searching
- Pattern matching
- Expression calculation

Transaction Management

Transaction Atomicity

- Transaction is a logical unit of such operations
- Each transaction finishes all operations or none

Transaction Management

Transaction Atomicity

- Transaction is a logical unit of such operations
- Each transaction finishes all operations or none

Locking Mechanisms

- Concurrency control mechanisms are required to protect data
- Two-Phase locking is implemented in Gromit

Transaction Management

Transaction Atomicity

- Transaction is a logical unit of such operations
- Each transaction finishes all operations or none

Locking Mechanisms

- Concurrency control mechanisms are required to protect data
- Two-Phase locking is implemented in Gromit


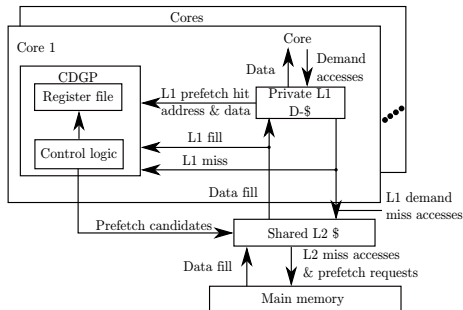
Deadlock Prevention Techniques

- No-Wait [5]
- Wait-Die
- Deadlock Detection

Benchmarks

Name	GDBench	LDBC-SNB
Vertex Types	2	13
Edge Types	2	20
Number of Queries	13	22
Query Example	Get the webpages liked by the friends of a person	Find a person's friends and friends of friends who started working in some company in a given country before a given year
Description	Data generator synthetically generates graphs that model social network activities with different connectivity.	LDBC-SNB models real-life social activities during a period of time.

Example usage of Gromit



1 Motivation

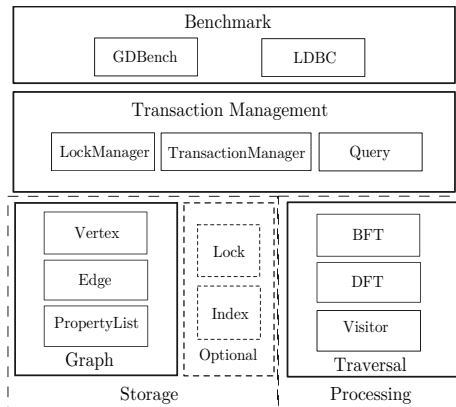
2 Main Contributions

3 Gromit

- Graph Storage
- Query Framework
- Transaction Management
- Benchmarks

4 Summary

Summary



<https://git.uwaterloo.ca/caesr-pub/gromit>

References



<http://cassandra.apache.org/>



<https://docs.mongodb.com/manual/>

▶ <http://all-free-download.com>



DeCandia et al. 2007. Amazon's Highly Available Key-value Store. *SIGOPS Oper. Syst. Rev.* 41, 6, 205-220.




Eswaran et al. 1976. The Notions of Consistency and Predicate Locks in a Database System. *Commun, ACM* 19, 11, 624-633.

Thank you

Question?

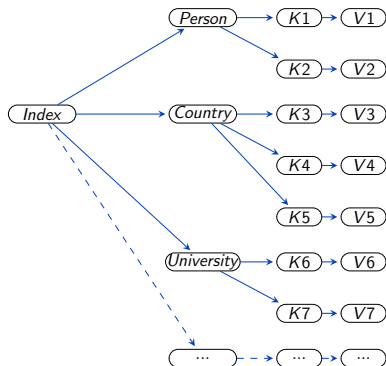
In-Memory or Disk



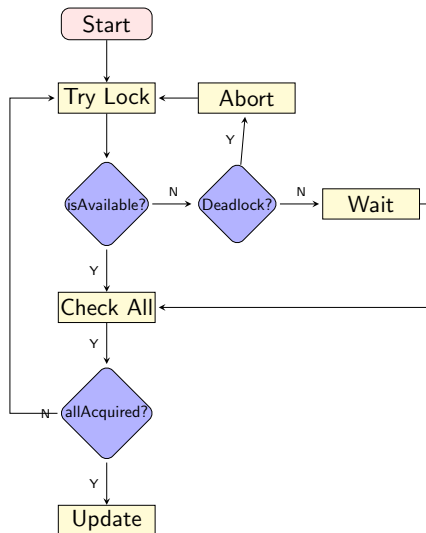
- Processor requests data from cache
- Data in cache are replaced by that from main memory
- Disk stores data and supplies to main memory
- Disk I/O can be bottleneck for memory-intensive workloads

Index

- Indexing retrieves information without traversing
- Indices are grouped by label
- Support for indexing is limited to unique keys in property list, such as ID



Deadlock Prevention Techniques



No-Wait

Never wait for a lock.
Abort right away.

Wait-Die

T_i waits for T_j only if $i < j$.
Otherwise, abort.

Detection

Construct *Wait-For Graph*
and check for cycles. If a
cycle exists, abort.
Otherwise, wait.